

# Getting Started with PostgreSQL

O'Reilly Online Live Training



I'm **Haki Benita**

A software developer and a technical lead.

I'm interested in databases, web development, software design and performance tuning.

**[hakibenita.com](https://hakibenita.com)**

# About Me

- Started as an Oracle DBA
- Worked on DWH and DSS systems
- Lead a team developers and DBAs
- Got into web development, mainly Python and Django
- Leading the development of a large ticketing platform

# POLL

## What is your job title?

- Developer
- Sysadmin / SRE
- Analyst / Data ops / Data scientist
- Student
- Management
- Other

# POLL

## What is your main use of data / SQL?

- Ad-Hoc Reports (sales reports, operational reports, etc.)
- Research / Analysis (BI / DWH / Dashboards/ Data exploration etc.)
- Development (Backend Development / ETL / Data pipelines etc.)

# POLL

**How would you rate your level of proficiency with databases?**

- Novice
- Intermediate
- Advanced

# The DBA Spectrum

## Infrastructure

- Operating system, storage, network, installation...
- Infra DBA, SRE, Sysadmin



# The DBA Spectrum

## Application

- SQL, ORMs, BI, dashboards, reporting...
- Developers, data scientists, analysts






# The DBA Spectrum

## Data

- Data model, tables, indexes, users, performance tuning
- Application DBAs, developers, data ops



# What you'll gain from this training

- Get comfortable with PostgreSQL
- Perform basic administrative tasks
  - Create and manage users and permissions
  - Create and manage tables
  - Evaluate query performance
  - Create indexes to speed up query execution

*Just the tip of the iceberg...*

*“ Give a man a fish and you feed him for a day  
Teach a man to fish and you feed him for a lifetime ”*

## In the process you will also

- Get comfortable with **PostgreSQL CLI and documentation**
- Figure out how to **find answers on your own**
- Evaluate different aspects of **database performance**


# History

# SQL

- **S**tructured **Q**uery **L**anguage
- Used for interacting with relational databases (RDBMS)
- Invented in the early 70s at IBM based on work by Edgar F. Codd
- Became a standard in 1986 (ANSI-86)
- Standard revised in 1992 (ANSI-92)

# PostgreSQL

- Based on the Berkeley POSTGRES project from 1986
- Originally named Postgres95
- Free Open Source
- Growing fast!




**Tools**



# PostgreSQL Documentation

- Excellent resource
- Suited for both experts and beginners
- Use it! Bookmark it!

<https://www.postgresql.org/docs/current/>



# psql

- PostgreSQL interactive terminal
- Comes with PostgreSQL installation

```
$ psql
psql (13.2)
Type "help" for help.

postgres=#
```

<https://www.postgresql.org/docs/13/app-psql.html>

# Don't memorize anything!

List all commands

```
postgres=# \?
```

```
General
```

```
\copyright          show PostgreSQL usage and distribution terms
\crosstabview [COLUMNS] execute query and display results in crosstab
\errverbose          show most recent error message at maximum verbosity
\g [(OPTIONS)] [FILE] e
....
```

# Don't memorize anything!

Get command syntax

```
postgres=# \h DELETE
```

```
Command:      DELETE
```

```
Description: delete rows of a table
```

```
Syntax:
```

```
[ WITH [ RECURSIVE ] with_query [, ...] ]
```

```
DELETE FROM [ ONLY ] table_name [ * ] [ [ AS ] alias ]
```

```
    [ USING from_item [, ...] ]
```

```
    [ WHERE condition | WHERE CURRENT OF cursor_name ]
```

```
    [ RETURNING * | output_expression [ [ AS ] output_name ] [, ...] ]
```

```
URL: https://www.postgresql.org/docs/13/sql-delete.html
```

# Don't memorize anything!

Use auto complete

```
postgres=# \set E<tab>  
ECHO          ECHO_HIDDEN          ENCODING
```

Also works with column and table names

**Architecture**

# Cluster


- PostgreSQL installation
- Define user / roles and their permissions
- Can contain multiple named databases

PostgreSQL Server 1



# Database


- Has a designated storage area on the file system
- Contains database objects and definitions such as tables, indexes and sequences
- Owned by a user






# User

- Granted object and system permissions
- Used to authenticate a client application



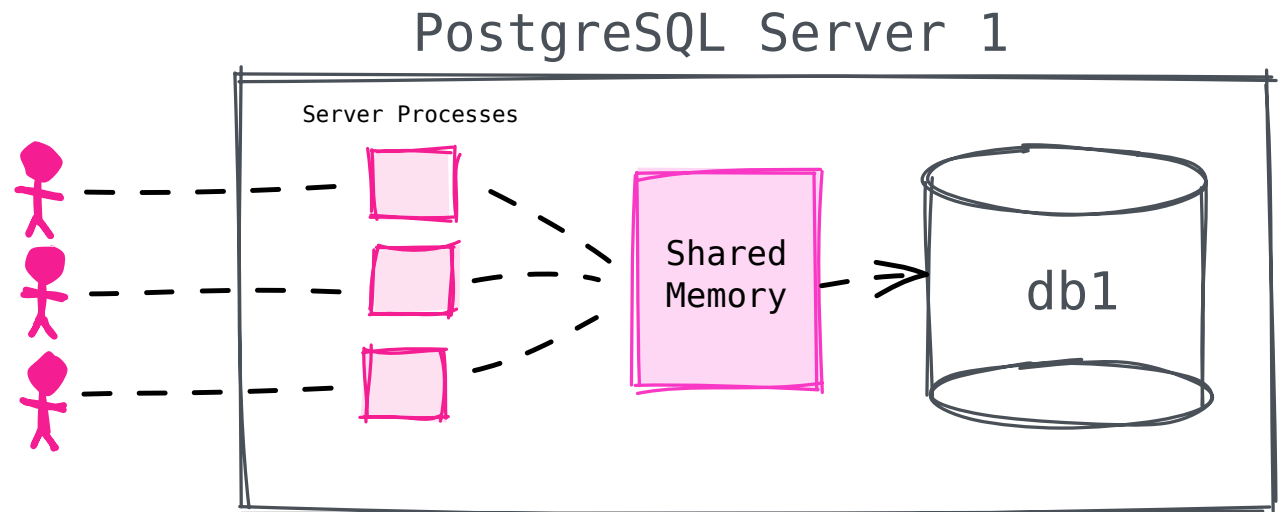
# Client Server

- Client application establish connection with database server
- Database server allocate a server process for the session
- Can have many concurrent active connections (  
`show max_connections`)



# Memory


- **Shared memory:** Shared by all server processes
- **Local memory:** Used by a single server process



# Memory

## Shared Buffers

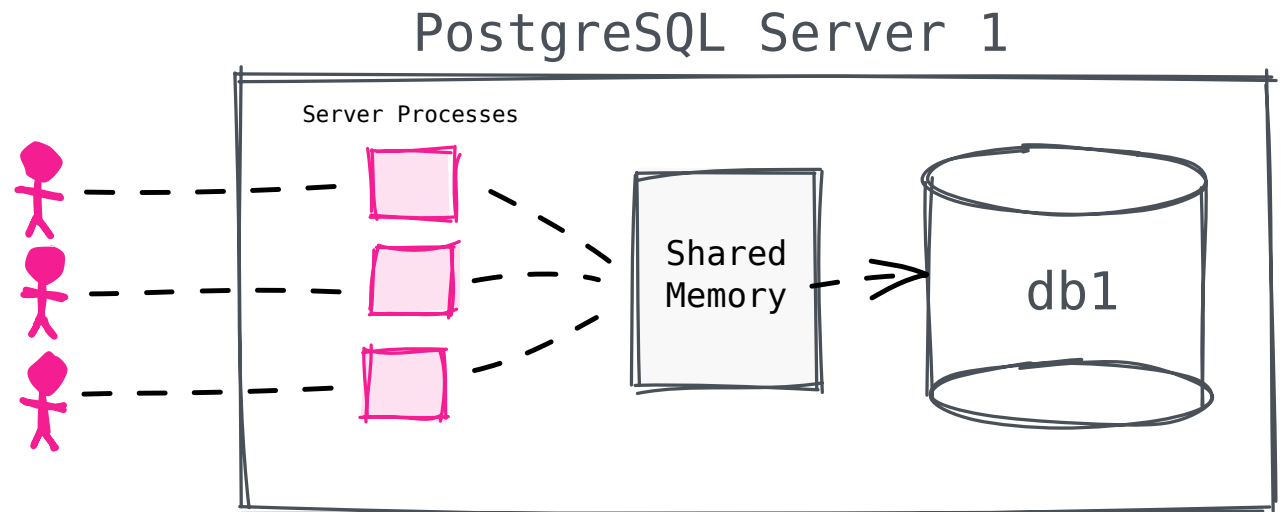
- Keep frequently accessed objects in memory for fast retrieval
- Shared by all server processes
- Default 128MB
- `show shared_buffers`



# Memory

## Work Mem


- Allocated for each backend process
- Used for sorts and hash tables
- Default 1MB
- `show work_mem`



# Memory

## Temp Buffers

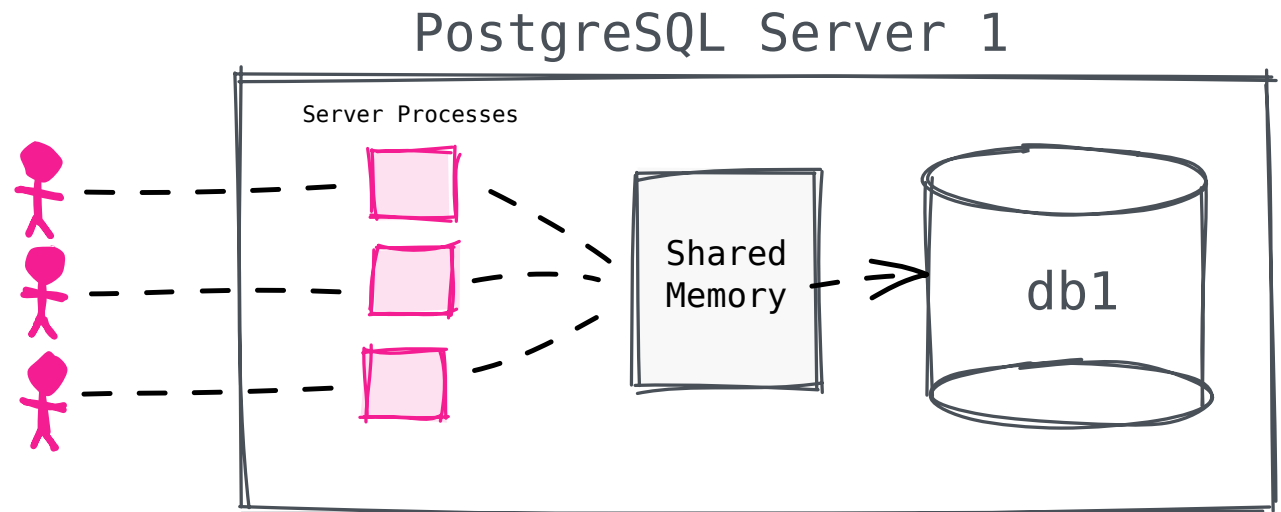
- Allocated for each backend process
- Used for storing temporary tables
- Default 8MB
- `show temp_buffers`



# Memory

## Maintenance Work Memory

- Allocated for each backend process
- Used for vacuum and create index operations
- Default 64MB
- `show maintenance_work_mem`



# Write Ahead Log (WAL)

- A log of all changes to tables and indexes
- Used to restore the database in case of disaster
- Can be used to maintain replication
- Enables point in time recovery




# Recap

- Architecture and Terminology
  - Cluster -> database -> schema -> table
- Memory structures
  - Shared memory
  - Local memory
- WAL

**Database**

# Database

- Contains database objects and definitions such as tables, indexes and sequences
- Owned by a user



# Database

## Creating a database

```
postgres=# CREATE DATABASE db1 OWNER postgres;  
CREATE DATABASE
```

```
postgres=# \connect db1  
You are now connected to database "db1" as user "postgres".  
db1=#
```

# Database

Creating a database from the terminal

```
$ createdb db1 -0 postgres
```

- `-0`: database owner

# Database

List databases

```
postgres=# \l
```

List of databases

Name	Owner	Encoding
db1	postgres	UTF8
postgres	postgres	UTF8
template0	postgres	UTF8
template1	postgres	UTF8

# Table

- Data in relational database is stored in tables
- Tables have columns
- Can define constraints
- Use indexes to speed access to data stored in tables

# Table

## Creating a table

```
db1=# CREATE TABLE users (  
    id INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
    active BOOLEAN,  
    name TEXT  
);  
  
CREATE TABLE
```



# Table

## Inspecting a table

```
db1=# \d users
```

Table "public.users"			
Column	Type	Nullable	Default
id	integer	not null	generated always as identity
active	boolean		
name	text		

Indexes:

```
"users_pkey" PRIMARY KEY, btree (id)
```

# Table

Listing all tables

```
db1=# \dt
```

List of relations

Schema	Name	Type	Owner
public	users	table	postgres

# Table

Searching for tables

```
db1=# \dt foo*
```

```
Did not find any relation named "foo*"
```

```
db1=# \dt u*
```

```
List of relations
```

Schema	Name	Type	Owner
public	users	table	postgres

# Table

Alter an existing table

```
db1=# ALTER TABLE users ALTER COLUMN active SET DEFAULT true;  
ALTER TABLE
```

```
db1=# \d users
```

Table "public.users"			
Column	Type	Nullable	Default
id	integer	not null	generated always as identity
active	boolean		true
name	text		

# Table

Default value is used when not explicitly provided

```
db1=# INSERT INTO users (name) VALUES ('Haki Benita');  
INSERT 0 1
```

```
db1=# SELECT * FROM users;
```

id	active	name
1	t	Haki Benita

Good defaults can prevent errors and confusion!

# View

- A named query
- Results are not materialized

# View

## Creating a view

```
db1=# CREATE VIEW active_users AS  
      SELECT id, name  
      FROM users  
      WHERE active;
```

```
CREATE VIEW
```

# View

## Querying a view

```
db1=# INSERT INTO users (name, active) VALUES ('Bob Bar', false);
INSERT 0 1
```

```
db1=# SELECT * FROM active_users;
```

id	name
1	Haki Benita

(1 row)

```
db1=# SELECT * FROM users;
```

id	active	name
1	t	Haki Benita
2	f	Bob Bar

(2 rows)



# View

## View details


db1=# \d+ active\_users

View "public.active_users"						
Column	Type	Collation	Nullable	Default	Storage	Description
id	integer				plain	
active	boolean				plain	
name	text				extended	

View definition:  
SELECT users.id, users.name  
FROM users  
WHERE users.active;

# Schema

- A namespace within the database
- Database can use multiple schemas
- The default schema is called "public"



# Schema

## Creating a schema

```
db1=# CREATE SCHEMA restricted;  
CREATE SCHEMA
```

```
db1=# \dn
```

List of schemas

Name	Owner
restricted	postgres
public	postgres

# Schema

## Creating tables in a schema

```
db1=# CREATE TABLE restricted.credentials (  
    id INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
    user_id INT,  
    password TEXT NOT NULL  
);  
CREATE TABLE
```

```
db1=# SELECT * FROM restricted.credentials;  
 id | user_id | password  
----|-----|-----  
(0 rows)
```

# Schema

Referencing objects by name

Name	Pattern	Example
qualified	<i>database.schema.table</i>	db1.restricted.credentials
qualified	<i>schema.table</i>	restricted.credentials
unqualified	<i>table</i>	credentials

# Schema

## Schema search path

```
db1=# SHOW search_path;  
      search_path  
-----  
"$user", public
```

- The schema order to look for unqualified object names
- `"$user"`: name of the current user
- `public`: the default schema

# Schema

## Schema search path

```
db1=# SELECT * FROM credentials;  
ERROR:  relation "credentials" does not exist
```

```
db1=# SET search_path TO restricted, "$user", public;  
SET
```

```
db1=# SELECT * FROM credentials;  
 id | user_id | password  
----+-----+-----  
(0 rows)
```

# Schema

Create objects with similar names in different schemas:


- Can act as a synonym
- Can be used for multi-tenancy, online data migrations

 ["Synonyms" in PostgreSQL](#)



# Information Schema

- A special schema in every database "pg\_catalog"
- Contains information about database objects



# Information Schema

Information about tables

```
db1=# SELECT * FROM pg_tables WHERE tablename = 'users';
-[ RECORD 1 ]-----
schemaname | public
tablename  | users
tableowner | postgres
tablespace | 
hasindexes | t
hasrules   | f
hastriggers | t
rowsecurity | f
```

# Information Schema

Information about indexes

```
db1=# SELECT * FROM pg_indexes WHERE tablename = 'users';
-[ RECORD 1 ]-----
schemaname | public
tablename  | users
indexname  | users_pkey
tablespace | 
indexdef   | CREATE UNIQUE INDEX users_pkey ON public.users USING btree (id)
```

# Information Schema

Information about all database objects:

```
db1=# SELECT relname, relkind FROM pg_class WHERE relname LIKE '%users%';
```

relname	relkind
users	r
users_pkey	i
users_id_seq	S
active_users	v

**r**=table, **i**=index, **s**=sequence, **v**=view, [more...](#)

# Information Schema

## Other useful tables

- `pg_stat_activity`: information about current activity
- `pg_stat_all_table`: table access statistics
- `pg_stat_all_indexes`: index access statistics
- `pg_stats`: statistics about table columns

# Recap

- Create database, schema, tables and views
- Referencing objects using qualified and unqualified names
- Schema search path
- Inspect and list database objects using psql
- Using the information schema (catalog)

# Exercise

# COPY

- Export data to text, CSV or binary format
- Import data from text, CSV or binary format into tables
- Produce reports from query results



# Exercise

Getting Started with PostgreSQL: Import and export data using COPY

**[Launch Katacoda scenario »](#)**

# Import From CSV

```
postgres=# \COPY users FROM /tmp/users.csv WITH CSV HEADER  
COPY 3
```

```
postgres=# \COPY users (active,name,id) FROM /tmp/users2.csv WITH CSV HEADER  
COPY 2
```

```
postgres=# SELECT * FROM users;
```

id	active	name
1	t	Haki Benita
2	f	Bob Bar
3	t	Ben Smith
10	t	Andres Lane
11	f	Jonathan Lane

# Export Table to CSV

```
postgres=# \COPY users TO /tmp/all_users.csv WITH CSV HEADER  
COPY 5
```

```
postgres=# \! cat /tmp/all_users.csv  
id,active,name  
1,t,Haki Benita  
2,f,Bob Bar  
3,t,Ben Smith  
10,t,Andres Lane  
11,f,Jonathan Lane
```

# Export Query Results to CSV

```
postgres=# \COPY (SELECT id, name FROM users WHERE active) TO /tmp/active_users.csv WITH CSV HEADER  
COPY 3
```

```
postgres=# \! cat /tmp/active_users.csv  
id,name  
1,Haki Benita  
3,Ben Smith  
10,Andres Lane
```

# COPY vs. \COPY

The COPY command has two variations:

- `COPY`: executed on the server
- `\COPY`: psql command with similar api, executed on the client

What you usually want is `\COPY`

# COPY Trick

You want to produce a report from a big query

```
SELECT u.id, u.name, count(*) AS credentials
FROM users u
  LEFT JOIN restricted.credentials c ON u.id = c.user_id
WHERE u.active
GROUP BY 1, 2
ORDER BY 3 DESC;
```

# COPY Trick

`\COPY` command does not support multiple lines

```
db1=# \copy (SELECT u.id, u.name, count(*) AS credentials <enter>
\copy: parse error at end of line
```

# COPY Trick

`COPY` can accept multiple lines

```
db1=# COPY (SELECT u.id, u.name, count(*) AS credentials
FROM users u LEFT JOIN restricted.credentials c ON u.id = c.user_id
WHERE u.active
GROUP BY 1, 2
ORDER BY 3 DESC) TO STDOUT WITH CSV HEADER;
id,name,credentials
1,Haki Benita,1
...
```

But you want to produce the report to a file, like `\copy` does...



# COPY Trick

`psql` has an option to send query output to file

```
db1=# \?  
...  
  \g [(OPTIONS)] [FILE]  
    execute query (and send results to file or |pipe);  
...
```

# COPY Trick

Use `\g` to write `COPY` output to a local file!

```
db1=# COPY (SELECT u.id, u.name, count(*) AS credentials
FROM users u LEFT JOIN restricted.credentials c ON u.id = c.user_id
WHERE u.active
GROUP BY 1, 2
ORDER BY 3 DESC)
TO STDOUT WITH CSV HEADER \g report.csv
```

```
COPY 4
```

 [Use \copy With Multi-line SQL](#)

# Data Integrity

Database as the source of truth

# Column Types

The first line of defence

```
db1=# INSERT INTO users (name, active) VALUES ('foo', 'not a boolean');  
ERROR:  invalid input syntax for type boolean: "not a boolean"
```

Cannot insert invalid values!

# Column Types

## Common types

- `varchar(N)`, `text`
- `smallint`, `integer`, `bigint`
- `decimal`, `real`, `double precision`
- `date`, `timestamp`, `timestampz`
- `bytea`
- `boolean`

# Column Types

## Special types

- uuid
- jsonb
- range
- array
- [more...](#)

# Column Types

Set restrictions that make sense

```
db1=# ALTER TABLE users ALTER name TYPE VARCHAR(20);  
ALTER TABLE
```

```
db1=# INSERT INTO users (name) VALUES ('probably invalid name');  
ERROR: value too long for type character varying(20)
```

 [Postgres: Boundless text and Back Again](#)

# Constraints

- Maintain data integrity
- Keeps the data clean
- Complements types



# **Not Null Constraint**

Make fields required

# Not Null Constraint

**NULL** is a special values that indicates "missing value"

# Not Null Constraint

Add a "not null" constraint on an existing table

```
db1=# ALTER TABLE users ALTER active SET NOT NULL;  
ALTER TABLE
```

```
db1=# \d users
```

Table "public.users"			
Column	Type	Nullable	Default
id	integer	not null	generated always as identity
active	boolean	not null	true
name	character varying(20)		

Indexes:

```
"users_pkey" PRIMARY KEY, btree (id)
```

# Not Null Constraint

Field is now required

```
db1=# INSERT INTO users (name, active) VALUES ('foo', NULL);  
ERROR: null value in column "active" of relation "users" violates not-null constraint  
DETAIL: Failing row contains (3, null, foo).
```

**active** must be set

# **Primary Key**

Primary unique identifier

# Primary Key

Set of fields that uniquely identify each row in the table

- Must be unique
- Must contain a value
- A table can only have one primary key
- A table does not have to have a primary key
- Can have multiple fields ("composite key")

# Primary Key

We already defined a primary key

```
db1=# CREATE TABLE users (  
    id INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
    active BOOLEAN,  
    name TEXT  
);
```

# Primary Key

We already defined a primary key

```
db1=# \d users
```

Table "public.users"			
Column	Type	Nullable	Default
id	integer	not null	generated always as identity
active	boolean		true
name	character varying(20)		

Indexes:

```
"users_pkey" PRIMARY KEY, btree (id)
```



# Primary Key

- Creates a unique index to enforce uniqueness
- Marks the fields as not null

# Auto Incrementing Primary Key

```
db1=# CREATE TABLE users (  
  id INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
  active BOOLEAN,  
  name TEXT  
);
```

- Creates a sequence
- Automatically populate PK with next values
- Used to be `serial` (soft-deprecated starting PostgreSQL 10)

# **Unique Constraint**

Ensure one or more fields are unique

# Unique Constraint

Name must be unique

```
db1=# ALTER TABLE users ADD CONSTRAINT users_name_unique UNIQUE(name);  
ALTER TABLE
```

# Unique Constraint

```
db1=# \d users
```

Table "public.users"				
Column	Type	Collation	Nullable	Default
id	integer		not null	generated always as identity
active	boolean		not null	true
name	character varying(20)			

Indexes:

```
"users_pkey" PRIMARY KEY, btree (id)
```

```
"users_name_unique" UNIQUE CONSTRAINT, btree (name)
```

Creates a unique index to enforce the constraint

# Unique Constraint

```
db1=# INSERT INTO users (name) VALUES ('Haki Benita');  
ERROR:  duplicate key value violates unique constraint "users_name_unique"  
DETAIL:  Key (name)=(Haki Benita) already exists.
```

Cannot insert duplicate values!

# Unique Constraint

Name is not really unique, so we can drop the constraint:

```
db1=# ALTER TABLE users DROP CONSTRAINT users_name_unique;  
ALTER TABLE
```

Will also drop the index

# **Check Constraint**

Implement special validation logic



# Check Constraint

Name must contain at least one whitespace

```
db1=# ALTER TABLE users ADD CONSTRAINT must_contain_whitespace CHECK (name LIKE '% %');  
ALTER TABLE
```

# Check Constraint

Name must contain at least one whitespace

```
db1=# \d users
```

Table "public.users"			
Column	Type	Nullable	Default
id	integer	not null	generated always as identity
active	boolean	not null	true
name	character varying(20)		

Indexes:

```
"users_pkey" PRIMARY KEY, btree (id)
```

Check constraints:

```
"must_contain_whitespace" CHECK (name::text ~~ '% '::text)
```

# Check Constraint

Name must contain at least one whitespace

```
db1=# INSERT INTO users (name) VALUES ('George');  
ERROR:  new row for relation "users" violates check constraint "must_contain_whitespace"  
DETAIL:  Failing row contains (4, t, George).
```

- Cannot set name without at least one whitespace
- Use meaningful names to make it easier to spot the problem

 [Add Constraints Without Validating Immediately](#)

# Check Constraint

💡 TIP: Use a check constraint with unrestricted text type if the length may change in the future

```
CREATE TABLE student (  
    id int,  
    name text,  
    CONSTRAINT name_length_check CHECK (length(name) < 20)  
);  
  
db1=# INSERT INTO student (id, name) VALUES (1, 'Haki Benita');  
INSERT 0 1  
  
db1=# INSERT INTO student (id, name) VALUES (2, 'Bernd Ottovordemgentschenfelde');  
ERROR:  new row for relation "student" violates check constraint "name_length_check"  
DETAIL:  Failing row contains (2, Bernd Ottovordemgentschenfelde).
```

# Check Constraint


💡 TIP: It is easier to change a check constraint than a column's data type

```
db1=# ALTER TABLE student DROP CONSTRAINT name_length_check;  
ALTER TABLE  
  
db1=# ALTER TABLE student ADD CONSTRAINT name_length_check CHECK(length(name) < 100);  
ALTER TABLE  
  
db1=# INSERT INTO student (id, name) VALUES (2, 'Bernd Ottovordemgentschenfelde');  
INSERT 0 1
```

📖 [Gitlab database guide: Strings and the Text data type](#)

# Foreign Key

Maintain relation between tables



# Foreign Key

Add a foreign key between users and their credentials

```
db1=# \d restricted.credentials
```

Column	Type
id	integer
user_id	integer
password	text

```
db1=# \d users
```

Column	Type
id	integer
active	boolean
name	character varying(20)

Indexes:

```
"users_pkey" PRIMARY KEY, btree (id)
```

# Foreign Key

Add a foreign key between users and their credentials

```
db1=# ALTER TABLE restricted.credentials ADD CONSTRAINT user_fk  
      FOREIGN KEY (user_id) REFERENCES users(id)  
      ON DELETE CASCADE;
```

```
ALTER TABLE
```



# Foreign Key

```
db1=# \d restricted.credentials
```

Table "restricted.credentials"			
Column	Type	Nullable	Default
id	integer	not null	generated always as identity
user_id	integer		
password	text	not null	

Indexes:

```
"credentials_pkey" PRIMARY KEY, btree (id)
```

Foreign-key constraints:

```
"user_fk" FOREIGN KEY (user_id) REFERENCES users(id) ON DELETE CASCADE
```

# Foreign Key

```
db1=# ALTER TABLE restricted.credentials ADD CONSTRAINT user_fk  
FOREIGN KEY (user_id) REFERENCES users(id)  
ON DELETE CASCADE;
```

```
ALTER TABLE
```

Column `user_id` must have a corresponding value in the `id` column in table `users`

# Foreign Key

Cannot add credentials for non-existing users

```
db1=# INSERT INTO restricted.credentials (user_id, password) VALUES (99, 'secret');  
ERROR:  insert or update on table "credentials" violates foreign key constraint "user_fk"  
DETAIL:  Key (user_id)=(99) is not present in table "users".
```

# Foreign Key

```
db1=# ALTER TABLE restricted.credentials ADD CONSTRAINT user_fk  
      FOREIGN KEY (user_id) REFERENCES users(id)  
      ON DELETE CASCADE;
```

ALTER TABLE

- **CASCADE**: When a user is deleted from the **users** table, delete the associated credentials in the **credentials** table
- **RESTRICT**: prevent deleting users with credentials

# Foreign Key

Add a new user with credentials

```
db1=# INSERT INTO users (name) VALUES ('Ben Smith') RETURNING *;
```

id	active	name
8	t	Ben Smith

```
INSERT 0 1
```

```
db1=# INSERT INTO restricted.credentials (user_id, password) VALUES (8, 'secret') RETURNING *;
```

id	user_id	password
2	8	secret

```
INSERT 0 1
```

# Foreign Key

Delete user

```
db1=# DELETE FROM users WHERE id = 8;  
DELETE 1
```

```
db1=# SELECT * FROM restricted.credentials WHERE user_id = 8;  
 id | user_id | password  
----|-----|-----  
(0 rows)
```

Credentials were also deleted

# Recap

- Make field required -> Not null constraint
- Custom validation logic -> Check constraint
- Ensuring uniqueness -> Unique constraint
- Maintain relation between tables -> Foreign key

# Conclusion

- Database is the source of truth
- Database is the last line of defence
- Data integrity is crucial


**Constraints are good... use constraints!**



# Exercise

# URL Shortener

- Create short URLs that redirect to longer URLs
- Keeps track of the number of clicks
- Common in SMS, Tweets
- Useful to keep track of campaigns



# Exercise

Getting Started with PostgreSQL: Creating and managing tables

**[Launch Katacoda scenario »](#)**

# **Users and Privileges**

# User

## Creating a user

```
db1=# CREATE USER app PASSWORD 'secret';  
CREATE ROLE
```

```
db1=# \connect db1 app  
You are now connected to database "db1" as user "app".
```

# User

## Connecting from terminal

```
$ psql -d db1 -U app -W  
Password:  
  
db1=>
```

- **-d** name of database
- **-U** name of user
- **-W** prompt for password

# Privileges

## Types of privileges

- SELECT / UPDATE / DELETE / TRUNCATE on TABLE
- CONNECT on DATABASE
- CREATE on DATABASE
- [more...](#)

# Privileges

Enforce privileges

```
db1=> \conninfo
You are connected to database "db1" as user "app".

db1=# SELECT * FROM users;
ERROR: permission denied for table users
```



# Privileges

Grant select privileges

```
db1=> \connect db1 postgres
You are now connected to database "db1" as user "postgres".

db1=# GRANT SELECT ON users TO app;
GRANT
```

# Privileges

Grant select privileges

```
db1=# \connect db1 app
```

```
You are now connected to database "db1" as user "app".
```

```
db1=> SELECT * FROM users;
```

id	active	name
1	t	Haki
2	f	Bob

# Privileges

Grant insert, update and delete privileges

```
db1=> \connect db1 postgres
You are now connected to database "db1" as user "postgres".

db1=# GRANT UPDATE, INSERT, DELETE ON users TO app;
GRANT
```

# Privileges

Use insert, update and delete privileges

```
db1=# \connect db1 app
```

```
You are now connected to database "db1" as user "app".
```

```
db1=> INSERT INTO users (active, name) VALUES (true, 'Jim');  
INSERT 0 1
```

```
db1=> UPDATE users SET active = false WHERE name = 'Jim';  
UPDATE 1
```

```
db1=> DELETE FROM users WHERE name = 'Jim';  
DELETE 1
```

# Privileges

## Revoke privileges

```
db1=> \connect db1 postgres
You are now connected to database "db1" as user "postgres".

db1=# REVOKE DELETE ON users FROM app;
REVOKE
```

Notice the user still has update, insert and select privileges

# Privileges

## Revoke privileges

```
db1=# \connect db1 app
You are now connected to database "db1" as user "app".
```

```
db1=> DELETE FROM users;
ERROR: permission denied for table users
```

# Privileges

Check user privileges on table

```
SELECT grantee, privilege_type
FROM information_schema.role_table_grants
WHERE table_name = 'users'
AND grantee = 'app';
```

grantee	privilege_type
app	INSERT
app	SELECT
app	UPDATE

# Privileges

Grant privileges on schema

```
db1=> \connect db1 postgres
You are now connected to database "db1" as user "postgres".

db1=> GRANT USAGE ON SCHEMA restricted TO app;
GRANT
```

Access to custom schemas require explicit **USAGE** privilege



# Privileges

Allow access to specific fields

```
db1=# GRANT SELECT (user_id) ON restricted.credentials TO app;  
GRANT
```

Useful for tables with sensitive information such as PII or passwords.

 [Grant Permissions on Specific Columns](#)

# Privileges

Allow access to specific fields

```
db1=> SELECT * FROM restricted.credentials;  
ERROR: permission denied for table credentials
```

```
db1=> SELECT user_id, password FROM restricted.credentials;  
ERROR: permission denied for table credentials
```

```
db1=> SELECT user_id FROM restricted.credentials;  
 user_id  
_____  
(0 rows)
```

# Roles

- Similar to user
- Don't have login privileges by default
- Should be thought of as a group of privileges a user can be a member of

# Roles

## Create role

```
db1=# \connect db1 postgres
You are now connected to database "db1" as user "postgres".
```

```
db1=# CREATE ROLE analyst;
CREATE ROLE
```

# Roles

Grant privileges to role

```
db1=# GRANT SELECT ON ALL TABLES IN SCHEMA public TO analyst;  
GRANT
```

```
db1=# GRANT UPDATE ON users TO analyst;  
GRANT
```

Syntax is similar to granting privileges to users

# Roles

Grant membership in a role:

```
db1=# CREATE USER bob;  
CREATE ROLE
```

```
db1=# GRANT analyst TO bob;  
GRANT
```

# Roles

User **bob** now has all the privileges the role **analyst** has

```
db1=# \connect db1 bob
```

```
You are now connected to database "db1" as user "bob".
```

```
db1=> SELECT * FROM users;
```

id	active	name
1	t	Haki
2	f	Bob

# Roles

List roles

```
db1=# \du
```

List of roles		
Role name	Attributes	Member of
analyst	Cannot login	{}
app		{}
bob		{analyst}
postgres	Superuser, Create role, Create DB, Replication, Bypass RLS	{}

- Notice **analyst** role cannot login
- Notice **bob** is member of **analyst**
- Notice **postgres** is a superuser



# Recap

- Create users and roles
- Grant and revoke privileges and roles
- Grant privileges on specific columns

# Query Optimizer

Where the magic happens...


# The Path of a Query

- **Parser:** Check for syntax errors
- **Rewrite:** Adjustments to the query (inline views etc.)
- **Planner / Optimizer:** Produce execution plan
- **Executer:** Execute the query according to the execution plan

# Query Optimizer

## 1. **Generate all possible execution plans**


- This can take some time, depending on the query
- [Cascade of doom: Postgres update that led to 70% failure](#)



# Query Optimizer

## 2. Estimate the cost for each plan


- Cost is measured in arbitrary units
- Using stats from analyzing tables and indexes
- Cost can be used to compare execution plans
- Mostly affected by IO



# Query Optimizer

## 3. Choose the plan with the lowest cost

- The lower the cost, the faster the execution is (expected) to be



# Execution Plan

Produce execution plan using the `EXPLAIN` command

```
db1=# EXPLAIN SELECT * FROM shorturl WHERE key = '123456';  
               QUERY PLAN
```

---

```
Index Scan using shorturl_key_key on shorturl (cost=0.29..8.30 rows=1 width=48)  
  Index Cond: ((key)::text = '123456'::text)
```

- Only produces an execution plan
- Does not execute the query

# Execution Plan

Produce execution plan and execute query

```
db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl WHERE key = '123456';
               QUERY PLAN
-----
Index Scan using shorturl_key_key on shorturl  (cost=0.29..8.30 rows=1 width=48)
                                                (actual time=0.032..0.033 rows=0 loops=1)
   Index Cond: ((key)::text = '123456'::text)
Planning Time: 0.121 ms
Execution Time: 0.064 ms
```

- Execute and time query
- Display executed plan
- Display estimated vs. actual estimates



# Execution Plan

## Reading an execution plan

```
db1=# EXPLAIN SELECT * FROM shorturl WHERE key = '123456';  
               QUERY PLAN
```

---

```
Index Scan using shorturl_key_key on shorturl (cost=0.29..8.30 rows=1 width=48)  
  Index Cond: ((key)::text = '123456'::text)
```

- The database used the index on **key** to find the row
- The optimizer estimates 1 row in the result

# Row Estimates

How the optimizer estimates how many rows in the result

# Row Estimates

Query all short URLs

```
db1=# EXPLAIN SELECT * FROM shorturl;  
               QUERY PLAN
```

---

```
Seq Scan on shorturl (cost=0.00..199.00 rows=10000 width=48)
```

- Database plans to scan the entire table
- Optimizer estimates 10,000 rows in the result

# Row Estimates

Using statistics on tables and indexes:

```
db1=# SELECT reltuples FROM pg_class WHERE relname = 'shorturl';  
reltuples  
-----  
10000
```

Table is *estimated* to have 10,000 tuples (rows)

# Row Estimates

Find short URLs with no hits:

```
db1=# EXPLAIN SELECT * FROM shorturl WHERE hits = 0;  
               QUERY PLAN
```

---

```
Seq Scan on shorturl (cost=0.00..224.00 rows=98 width=48)  
  Filter: (hits = 0)
```

- Database plans to scan the entire table
- Optimizer estimates 98 rows in the result

# Row Estimates

Using statistics on table columns:

```
db1=# SELECT tablename, attname, most_common_vals, most_common_freqs
FROM pg_stats
WHERE tablename = 'shorturl' AND attname = 'hits';
```

```
—[ RECORD 1 ]———  
tablename      | shorturl  
attname         | hits  
most_common_vals | {0, 7306, 9658, ...  
most_common_freqs | {0.0098, 0.0008, 0.0007, ...
```

number of rows (10,000) \* % hits eq 0 (0.0098) = 98 rows

# Statistics

- The database keeps statistics on tables, indexes and columns
- Statistics can be collected explicitly using **ANALYZE**
- The database can collect statistics in the background

# Statistics

Explicitly collect statistics on a table:

```
db1=# ANALYZE shorturl;  
ANALYZE
```



# Statistics

Collect statistics automatically in the background:

```
db1=# show autovacuum;  
autovacuum  
-----  
on
```

- "autovacuum" is a routine maintenance task performed by the database
- The database can analyze tables and indexes as they are vacuumed

# Statistics

Check when analyzed:

```
db1=# SELECT last_analyze, last_autoanalyze
FROM pg_stat_all_tables
WHERE relname = 'shorturl';
```

```
-[ RECORD 1 ]-----
last_analyze   | 2021-08-05 11:09:45.228915+03
last_autoanalyze | 2021-08-05 09:58:40.550875+03
```

# Recap

- Optimizer uses statistics to produce execution plans
- Keeping statistics up to date has a significant effect on query performance
- Statistics can be collected explicitly using `ANALYZE`, or in the background by enabling autovacuum

# Indexes

The need for speed

# Indexes

- Speed us access to data
- Enforce constraints (Unique, Foreign Key)
- PostgreSQL offers many different types of indexes

# **B-Tree Index**

The king of all indexes!

# Exercise

# Exercise

## Getting Started with PostgreSQL: B-Tree Index Features

- Partial B-Tree index
- Inclusive B-Tree index
- Function based B-Tree index

**[Launch Katacoda scenario »](#)**




# B-Tree Index

How does a B-Tree index work?

1. You have these values:

1 2 3 4 5 6 7 8 9


2. Create a tree



# B-Tree Index

Search for the value **5**:

1. Scan the index root
2. Find the leaf block
3. Search the value in the leaf



# B-Tree Index

Create a B-Tree index

```
db1=# CREATE INDEX shorturl_hits_ix ON shorturl USING btree(hits);  
CREATE INDEX
```

Can omit `USING btree()`, this is the default

# Partial Index

Index only a portion of the table

```
db1=# CREATE INDEX shorturl_unused_part_ix ON shorturl (id) WHERE hits = 0;  
CREATE INDEX
```

Index only short URLs with zero hits

# Partial Index

Smaller index

```
db1=# \di+ shorturl*_ix
```

Name	Type	Table	Size
-----+-----+-----+			
shorturl_hits_ix	index	shorturl	240 kB
shorturl_unused_part_ix	index	shorturl	16 kB

# Partial Index

- Produce smaller indexes
- Limited to queries using the indexed rows
- Nullable columns are great candidates
- Use when possible

 [The Unexpected Find That Freed 20GB of Unused Index Space](#)

# Inclusive Index

Store additional data in the index leafs:

```
db1=# CREATE UNIQUE INDEX shorturl_key_including_url_ix ON shorturl(key) INCLUDE (url);  
CREATE INDEX
```

Index **key** and include the value of **url** in the index leaf

# Inclusive Index

Fulfill queries using just the index:

```
db1=# EXPLAIN (ANALYZE, TIMING) SELECT url FROM shorturl WHERE key = 'key123';  
QUERY PLAN
```

```
-----  
Index Only Scan using shorturl_key_including_url_ix on shorturl
```

```
Index Cond: (key = 'key123'::text)
```

```
Heap Fetches: 0
```

```
Planning Time: 0.534 ms
```

```
Execution Time: 0.120 ms
```

Notice "Index Only Scan"



# Inclusive Index

- Fulfill queries without accessing the table
- Index can get very big
- Great for large tables with frequent queries that only use a limited number of columns
- Use with caution
- Non-unique composite indexes can be good candidates for inclusive indexes

# Function Based Index

Index an expression

```
db1=# CREATE INDEX shorturl_domain_ix ON shorturl (substring(url FROM '.*://([^\s/]+)'));  
CREATE INDEX
```

Index domain part of URL using regular expression (regexp)

<https://hakibenita.com/sql-for-data-analysis>

# Function Based Index

Can be used by queries using the same expression

```
-- No index
db1=# SELECT * FROM shorturl WHERE substring(url FROM '.*://([^/]+)') = 'hakibenita.com';
Execution Time: 48.918 ms

-- With index
db1=# SELECT * FROM shorturl WHERE substring(url FROM '.*://([^/]+)') = 'hakibenita.com';
Execution Time: 0.914
```

Much faster!

# Function Based Index

- Index will only be considered only if the expressions is exactly the same
- Useful for existing applications (you can't change)
- Consider calculated columns as an alternative

# Recap

- B-Tree is the default index type and what you normally need
- Used to enforce unique constraints
- Many features:
  - Partial
  - Function based
  - Inclusive
  - [more...](#)

# **Performance**

What's it all about

# Performance

- CPU
- Memory
- Disk Space
- Cost

Not just speed!

# Table Size

Functions to get table size

Function	Description
<code>pg_relation_size</code>	Table size
<code>pg_table_size</code>	Including TOAST
<code>pg_total_relation_size</code>	Including TOAST & Indexes

TOAST is extended storage for large values

 [The Surprising Impact of Medium-Size Texts on PostgreSQL Performance](#)



# Table Size

```
db1=# SELECT pg_table_size('shorturl');  
pg_table_size
```

---

```
851968
```

```
db1=# SELECT pg_size_pretty(pg_table_size('shorturl'));  
pg_size_pretty
```

---

```
832 kB
```

```
db1=# \dt+ shorturl
```

List of relations					
Schema	Name	Type	Owner	Persistence	Size
public	shorturl	table	haki	permanent	832 kB

# Index Size

Check index size

```
db1=# \di+ shorturl*
```

List of relations

Schema	Name	Type	Owner	Table	Persistence	Size
public	shorturl_key_key	index	haki	shorturl	permanent	312 kB
public	shorturl_pkey	index	haki	shorturl	permanent	240 kB

Naming convention makes this easy...

# Hash Index

The Ugly Duckling of index types

# Reverse Lookup

Find keys referencing a URL

```
SELECT *  
FROM shorturl  
WHERE url = 'https://hakibenita.com/postgresql-hash-index';
```

- URL is not unique, but it is *almost unique*
- URL can be a large text

# Reverse Lookup

No index

```
db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl
WHERE url = 'https://hakibenita.com/postgresql-hash-index';
```

## QUERY PLAN

```
-----
Seq Scan on shorturl (cost=0.00..2239.00 rows=1 width=48)
  Filter: (url = 'https://hakibenita.com/postgresql-hash-index'::text)
  Rows Removed by Filter: 99998
Planning Time: 0.139 ms
Execution Time: 15.506 ms
```

Full table scan

# Reverse Lookup

## B-Tree index

```
db1=# CREATE INDEX shorturl_url_ix ON shorturl (url);  
CREATE INDEX  
db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl  
WHERE url = 'https://hakibenita.com/postgresql-hash-index';
```

### QUERY PLAN

```
-----  
Index Scan using shorturl_url_ix on shorturl (cost=0.42..8.44 rows=1 width=48)  
  Index Cond: (url = 'https://hakibenita.com/postgresql-hash-index'::text)  
Planning Time: 0.334 ms  
Execution Time: 0.109 ms
```

## Index scan using B-Tree index

# Reverse Lookup

Recap

Access method	Reverse lookup timing
Full table scan	15.506 ms
B-Tree index scan	0.109 ms

# Hash Index

How does a Hash index work?

Value	
A	
B	
C	
D	



# Hash Index

Apply a hash function on the values ( `hashchar` )

Value	<code>hashchar(value)</code>
A	-201530951
B	626080936
C	2018813598
D	2016322020

Hash functions for other types

`hashtext`

`hashchar`

`hash_array`

`jsonb_hash`

`timestamp_hash`

# Hash Index

Divide to N buckets using `mod(N)` :

Value	hashchar(value)	Bucket mod(2)
A	-201530951	1
B	626080936	0
C	2018813598	0
D	2016322020	0

# Hash Index

Build the index:

Bucket	Row Pointers
0	2, 3, 4
1	1

# Hash Index

Search for the value **B**:

1. Find bucket

```
hashchar('B') -> 626080936
```

```
mod(626080936, 2) -> 0
```

2. Scan rows in bucket 0 and search value:

Bucket	Row Pointers
0	2, 3, 4

# Reverse Lookup

## Hash index

```
db1=# CREATE INDEX shorturl_url_hix ON shorturl USING HASH (url);  
CREATE INDEX  
db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl  
WHERE url = 'https://hakibenita.com/postgresql-hash-index';
```

### QUERY PLAN

```
-----  
Index Scan using shorturl_url_hix on shorturl (cost=0.00..8.02 rows=1 width=48)  
  Index Cond: (url = 'https://hakibenita.com/postgresql-hash-index'::text)  
Planning Time: 0.242 ms  
Execution Time: 0.061 ms
```

## Index scan using Hash index

# Reverse Lookup

B-Tree vs. Hash Index speed

Access method	Reverse lookup timing
Full table scan	15.506 ms
B-Tree index scan	0.109 ms
Hash index scan	0.061 ms

Hash index is faster!

# Reverse Lookup

## B-Tree vs. Hash Index size

```
db1=# \di+ shorturl_url_*
```

List of relations


Schema	Name	Type	Owner	Table	Size
public	shorturl_url_hix	index	postgres	shorturl	4112 kB
public	shorturl_url_ix	index	postgres	shorturl	4584 kB

Hash index is smaller!

# Hash Index

## Hash vs. B-Tree index size

Hash vs. B-Tree Index Size





# Hash Index Restrictions



- ✗ Unique
- ✗ Composite
- ✗ Sorting
- ✗ Range search

# Hash Index

- Ideal when values are *almost unique*
- Not affected by the size of the values
- Can be smaller and faster than a B-Tree
- Was discouraged prior to PostgreSQL 10, but no more!

# Hash Index

Go Further:

-  [Getting Started with PostgreSQL: Hash Index](#)  
Interactive Katacode scenario
-  [Re-Introducing Hash Indexes in PostgreSQL](#)

# Block Range Index

Keep range of values within a number of adjacent pages

# Range Search

Find short URLs that were created in a date range

```
SELECT *  
FROM shorturl  
WHERE created_at >= '2021-02-01 UTC'  
AND created_at < '2021-03-01 UTC';
```

- Creation date set by the application when a row is added
- Creation date is naturally incrementing

# Range Search

No Index

```
db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl
WHERE created_at >= '2021-02-01 UTC' AND created_at < '2021-03-01 UTC';
               QUERY PLAN
```

```
-----
Seq Scan on shorturl
```

```
  Filter: ((created_at >= '2021-02-01 00:00:00+00'::timestamp with time zone)
    AND (created_at < '2021-03-01 00:00:00+00'::timestamp with time zone))
```

```
  Rows Removed by Filter: 95967
```

```
Execution Time: 19.003 ms
```

Full table scan

# Range Search

## B-Tree index

```
db1=# CREATE INDEX shorturl_created_at_ix ON shorturl (created_at);
CREATE INDEX

db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl
WHERE created_at >= '2021-02-01 UTC' AND created_at < '2021-03-01 UTC';
QUERY PLAN

-----
Index Scan using shorturl_created_at_ix on shorturl
  Index Cond: ((created_at >= '2021-02-01 00:00:00+00'::timestamp with time zone)
    AND (created_at < '2021-03-01 00:00:00+00'::timestamp with time zone))
Execution Time: 2.178 ms
```

## B-Tree Index scan

# Range Search

Recap

Access method	Index Size	Timing
Full table scan	-	19.003 ms
B-Tree index scan	2208 kB	2.178 ms



# BRIN Index

How does a BRIN index work?

1. You have these values in a column, each is single table page:

1 2 3 4 5 6 7 8 9

# BRIN Index

How does a BRIN index work?

1. You have these values in a column, each is single table page:

1 2 3 4 5 6 7 8 9

2. Divide the table into ranges of 3 adjacent pages:

[1,2,3] [4,5,6] [7,8,9]

# BRIN Index

How does a BRIN index work?

1. You have these values in a column, each is single table page:

1 2 3 4 5 6 7 8 9

2. Divide the table into ranges of 3 adjacent pages:

[1,2,3] [4,5,6] [7,8,9]

3. For each range, keep only the minimum and maximum values:

[1-3] [4-6] [7-9]

# BRIN Index

Use the index to search for the value 5:

- [1–3] 🦴 Definitely not here
- [4–6] 💩 Might be here
- [7–9] 🦴 Definitely not here

We only need to scan blocks 4, 5 and 6!

# Range Search

## BRIN index

```
db1=# CREATE INDEX shorturl_created_at_ix ON shorturl USING brin(created_at);
CREATE INDEX

db1=# EXPLAIN (ANALYZE, TIMING) SELECT * FROM shorturl
WHERE created_at >= '2021-02-01 UTC' AND created_at < '2021-03-01 UTC';
QUERY PLAN

-----
Bitmap Heap Scan on shorturl
  Recheck Cond: ((created_at >= '2021-02-01 00:00:00+00'::timestamp with time zone) AND ...)
  Rows Removed by Index Recheck: 8921 Heap Blocks: lossy=128
-> Bitmap Index Scan on shorturl_created_at_bix
    Index Cond: ((created_at >= '2021-02-01 00:00:00+00'::timestamp with time zone) AND ...)
Execution Time: 5.207 ms
```

## Bitmap Index Scan on BRIN index

# Range Search

BRIN index size

```
db1=# \di+ shorturl_created_at_*
```

Name	Type	Table	Size
-----+-----+-----+			
shorturl_created_at_ix	index	shorturl	2208 kB
shorturl_created_at_bix	index	shorturl	48 kB

BRIN is very small!

# Range Search

BRIN vs. B-Tree

Access method	Index Size	Timing
Full table scan	-	19.003 ms
B-Tree index scan	2208 kB	<b>2.178 ms</b>
BRIN index scan	<b>48 kB</b>	5.207 ms

- B-Tree is faster
- BRIN is smaller

# BRIN Index

What if the values are **not sorted**?

1. 2 9 5 1 4 7 3 8 6
2. [2, 9, 5] [1, 4, 7] [3, 8, 6]
3. [2-9] [1-7] [3-8]



# BRIN Index

Use the index to search for the value 5:

- [2-9] 🤩 Might be here
- [1-7] 🤩 Might be here
- [3-8] 🤩 Might be here

The index is useless!

# Correlation

Correlation between logical and physical ordering:

```
db1=# SELECT attname, correlation FROM pg_stats WHERE tablename = 'shorturl';
```

attname	correlation
id	1
key	0.0018741399
url	0.0033386275
hits	0.0015433382
created_at	1

1=incrementing -1=decreasing ~0=not correlated

# BRIN Index

What if we increase the number of **pages per range** to 5?

1. **1** **2** **3** **4** **5** **6** **7** **8** **9**

2. **[1,2,3,4,5]** **[6,7,8,9]**

3. **[1-5]** **[6,9]**

# BRIN Index

Use the index to search for the value 5:

- [1-5] 🤔 Might be here
- [6-9] 💀 Definitely not here

We need to scan blocks 1, 2, 3, 4 and 5

# BRIN Index

What if we decrease the number of **pages per range** to 2?

1. **1** **2** **3** **4** **5** **6** **7** **8** **9**
2. **[1,2]** **[3,4]** **[5,6]** **[7,8]** **[9]**
3. **[1-2]** **[3-4]** **[5-6]** **[7-8]** **[9-9]**

# BRIN Index

Use the index to search for the value 5:

- [1-2] 🦴 Definitely not here
- [3-4] 🦴 Definitely not here
- [5-6] 💩 Might be here
- [7-8] 🦴 Definitely not here
- [9-9] 🦴 Definitely not here

We only need to scan blocks 5 and 6

# BRIN Index: Pages per range

```
CREATE INDEX shorturl_created_at_ix ON shorturl  
USING brin(created_at) WITH (pages_per_range = 128);
```

- Low `pages_per_range` -> more accurate, bigger size
- High `pages_per_range` -> less accurate, smaller size
- Default is 128, minimum is 2
- Start small!





# BRIN Index

- Ideal when data is naturally sorted on disk
  - Find columns with high correlation in `pg_stats.correlation`
  - Columns for auto incrementing column (timestamps etc.)
- Ideal for tables that don't update frequently
- Adjust `pages_per_range` to find ideal range size



# BRIN Index

Go Further:

-  [Getting Started with PostgreSQL: Block range index](#)  
Interactive Katacode scenario
-  [Index Columns With High Correlation Using BRIN](#)
-  [9 Django Tips for Working with Databases: BRIN indexes](#)
-  [Multi minmax and Bloom operators for BRIN index](#)  
New in PostgreSQL 14

# Conclusion

# You learned

- Architecture and terminology
- Manage tables and indexes
- Create constraints and choose appropriate data types
- Manage users, roles and privileges
- How the database produces execution plans
- How to create different types of indexes

## **You *also* learned**

- How to be productive using psql
- Maintain data integrity using data types and constraints
- Evaluate different aspects of query performance
- Optimize query performance using indexes

# **See you next time!**

[SQL Next Steps: Optimization](#)

Online Live Training, Aug 2, 2022

# Haki Benita



Check out my blog [hakibenita.com](https://hakibenita.com)



Find me on Twitter [@be\\_haki](https://twitter.com/be_haki)



Subscribe to my newsletter at [hakibenita.com/subscribe](https://hakibenita.com/subscribe)



Send me an email [me@hakibenita.com](mailto:me@hakibenita.com)